

Fundamentals of Data Science – Simple Python Programs

Name : M.Greeshma

Reg No : 192524195

Code : DSA0409

Slot : A

Question 1: Student Subject Averages using NumPy

```
[6]
✓ Os
import numpy as np
student_scores = np.array([
    [80, 75, 85, 70],
    [90, 85, 80, 75],
    [70, 80, 90, 85],
    [85, 90, 95, 80]
])
subjects = ["Math", "Science", "English", "History"]
average = np.mean(student_scores, axis=0)
for i in range(4):
    print(subjects[i], "Average =", average[i])
highest = np.argmax(average)
print("Highest average subject:", subjects[highest])

...
Math Average = 81.25
Science Average = 82.5
English Average = 87.5
History Average = 77.5
Highest average subject: English
```

Question 2: Average Product Price using NumPy

```
[4]
✓ Os
import numpy as np
price=np.array([
    [10,20,30],
    [40,50,60],
    [70,80,90]
])
average_prices=np.mean(prices)
print("average prices:",average_prices)

average prices: 130.0
```

Question 3: Average Sale Price of Houses with More Than 4 Bedrooms

```
[7]
✓ 0s
import numpy as np
# Columns: Bedrooms, Area, Sale Price
house_data = np.array([
    [3, 1200, 250000],
    [5, 1800, 400000],
    [6, 2200, 500000],
    [4, 1500, 320000],
    [5, 2000, 450000]
])
selected = house_data[house_data[:, 0] > 4]
average_price = np.mean(selected[:, 2])
print("Average sale price =", average_price)

... Average sale price = 450000.0
```

Question 4: Total Yearly Sales and Percentage Increase

```
[8]
✓ 0s
import numpy as np

sales_data = np.array([10000, 12000, 15000, 18000])

total_sales = np.sum(sales_data)
percentage_increase = ((int64: total_sales - data[0]) /
                       np.int64(55000))

print("Total sales =", total_sales)
print("Percentage increase =", percentage_increase, "%")

... Total sales = 55000
   Percentage increase = 80.0 %
```

Question 5: Average Fuel Efficiency and Percentage Improvement

```
[9]
✓ 0s
import numpy as np

fuel_efficiency = np.array([25, 30, 28, 35, 32])

average = np.mean(fuel_efficiency)

old_model = fuel_efficiency[0]
new_model = fuel_efficiency[1]

improvement = ((new_model - old_model) / old_model) * 100

print("Average fuel efficiency =", average)
print("Percentage improvement =", improvement, "%")

... Average fuel efficiency = 30.0
   Percentage improvement = 20.0 %
```

Question 6: Total Grocery Cost with Discount and Tax

```
[10]
✓ 0s
▶ prices = [100, 50, 200]
  quantities = [2, 3, 1]

  discount_rate = 10
  tax_rate = 5

  subtotal = 0

  for price, quantity in zip(prices, quantities):
      subtotal += price * quantity

  discount = subtotal * discount_rate / 100
  after_discount = subtotal - discount

  tax = after_discount * tax_rate / 100
  total = after_discount + tax

  print("Subtotal =", subtotal)
  print("Discount =", discount)
  print("Tax =", tax)
  print("Total cost =", total)

▼ ... Subtotal = 550
  Discount = 55.0
  Tax = 24.75
  Total cost = 519.75
```

Question 7: Pandas Order Analysis

```
[12]
✓ 0s
▶ print("Orders by customer:")
  print(data.groupby("Customer_ID").size())

  print("\nAverage quantity by product:")
  print(data.groupby("Product")["Quantity"].mean())

  data["Order_Date"] = pd.to_datetime(data["Order_Date"])

  print("\nEarliest date:", data["Order_Date"].min())
  print("Latest date:", data["Order_Date"].max())

▼ ... Orders by customer:
  Customer_ID
1      2
2      2
3      1
dtype: int64

  Average quantity by product:
  Product
Bag      1.0
Book     3.5
Pen      3.5
Name: Quantity, dtype: float64

  Earliest date: 2026-01-02 00:00:00
  Latest date: 2026-01-10 00:00:00
```

Question 8: Top 5 Products Sold

```
[13]
✓ Os  import pandas as pd

sales = pd.DataFrame({
    "Product": ["Pen", "Book", "Pen", "Bag", "Book", "Pen", "Bag"],
    "Quantity": [10, 8, 15, 5, 12, 20, 9]
})

top5 = sales.groupby("Product")["Quantity"].sum()
top5 = top5.sort_values(ascending=False).head(5)

print(top5)
```

```
✓ ... Product
    Pen    45
    Book   20
    Bag    14
Name: Quantity, dtype: int64
```

Question 9: Property Data Analysis using Pandas

```
[14]
✓ Os  "Area": [1200, 1800, 1500, 2200, 2000],
    "Price": [300000, 450000, 350000, 500000, 480000]
})


print("Average price by location:")
print(property_data.groupby("Location")["Price"].mean())

print("Properties with more than 4 bedrooms:")
print((property_data["Bedrooms"] > 4).sum())

largest = property_data.loc[property_data["Area"].idxmax()]
print("Property with largest area:")
print(largest)
```



```
✓ ... Average price by location:
Location
Chennai    410000.0
Madurai    425000.0
Name: Price, dtype: float64
Properties with more than 4 bedrooms:
3
Property with largest area:
Location    Madurai
Bedrooms      6
Area        2200
Price      500000
Name: 3, dtype: object
```


Question 10: Monthly Sales – Line Plot and Bar Plot




Question 11: Sales over Time – Line, Scatter and Bar Plots


```
[16] plt.ylabel("Sales")  
plt.show()
```





Question 12: Monthly Temperature and Rainfall Plots





Question 13: Word Frequency from a Text File

```
[21] ✓ 0s from collections import Counter
import re


text = "Python is easy. Python is powerful. Python is easy."

words = re.findall(r"\w+", text.lower())
count = Counter(words)


for word, value in count.items():
    print(word, ":", value)
```

```
python : 3
is : 3
easy : 2
powerful : 1
```

Question 14: Frequency Distribution of Customer Ages



Question 15: Frequency Distribution of Likes



Question 16: Word Frequency in Customer Reviews

```
from collections import Counter

reviews = pd.DataFrame({
    "Review": [
        "good product and good service",
        "excellent product",
        "good quality product"
    ]
})

words = []


for review in reviews["Review"]:
    words.extend(re.findall(r"\b\w+\b", review.lower()))

frequency = Counter(words)


for word, count in frequency.most_common():
    print(word, ":", count)
```


... good : 3
product : 3
and : 1
service : 1
excellent : 1
quality : 1

Question 17: Customer Feedback – Top N Word Frequency



Question 18: Mean, Median, Standard Deviation, Boxplot, Scatter Plot and Q-Q Plot





Question 19: 95% Confidence Interval for Blood Pressure Reduction

```
[28]
✓ Os  import numpy as np
    from scipy import stats

    drug = np.array([8, 10, 12, 9, 11, 13, 10, 12])
    placebo = np.array([3, 5, 4, 6, 2, 5, 4, 3])

    def confidence_interval(data):
        mean = np.mean(data)
        sem = stats.sem(data)
        interval = stats.t.interval(
            0.95, len(data) - 1, loc=mean, scale=sem
        )
        return interval

    print("Drug 95% CI:", confidence_interval(drug))
    print("Placebo 95% CI:", confidence_interval(placebo))
```

✓ ... Drug 95% CI: (np.float64(9.216289685290807), np.float64(12.033710314709193))
Placebo 95% CI: (np.float64(2.9053916697328868), np.float64(5.094608330267113))

Question 20: A/B Test – Compare Two Conversion Rate Samples

```
[29]
✓ 0s
import numpy as np
from scipy.stats import ttest_ind

design_A = np.array([0.10, 0.12, 0.11, 0.13, 0.09, 0.12])
design_B = np.array([0.14, 0.15, 0.13, 0.16, 0.12, 0.14])

t_stat, p_value = ttest_ind(design_A, design_B)

print("t-statistic =", t_stat)
print("p-value =", p_value)

if p_value < 0.05:
    print("Statistically significant difference")
else:
    print("No statistically significant difference")

... t-statistic = -3.3999999999999995
p-value = 0.006771072424337956
Statistically significant difference
```

Question 21: Point Estimate and Confidence Interval from CSV

```
[31]
✓ 26s
n = int(input("Enter sample size: "))
confidence = float(input("Enter confidence level (example 0.95): "))
precision = float(input("Enter desired precision: "))

sample = values.sample(n=min(n, len(values)), random_state=1)

mean = sample.mean()
sem = stats.sem(sample)

ci = stats.t.interval(
    confidence,
    len(sample) - 1,
    loc=mean,
    scale=sem
)

print("Point estimate =", mean)
print("Confidence interval =", ci)
print("Desired precision =", precision)
|

... Enter sample size: 5
Enter confidence level (example 0.95): 0.95
Enter desired precision: 13
Point estimate = 23.0
Confidence interval = (np.float64(11.96384898298353), np.float64(34.03615101701647))
Desired precision = 13.0
```

Question 22: Confidence Interval for Mean Customer Rating

```
[34] import pandas as pd
      from scipy import stats

      data = pd.DataFrame({
          "rating": [4, 5, 3, 4, 5, 4, 5, 3, 4, 5]
      })

      ratings = data["rating"].dropna()


      mean = ratings.mean()
      sem = stats.sem(ratings)

      ci = stats.t.interval(
          0.95,
          len(ratings) - 1,
          loc=mean,
          scale=sem
      )

      print("Mean rating =", mean)
      print("95% Confidence interval =", ci)
```

... Mean rating = 4.2
95% Confidence interval = (np.float64(3.635718862770892), np.float64(4.7642811372291085))

Question 23: Hypothesis Testing for Treatment vs Placebo



Question 24: KNN Classifier for a New Patient

```
[36] ✓ 1s from sklearn.neighbors import KNeighborsClassifier

# Features: symptom1, symptom2
X = [
    [1, 1], [1, 2], [2, 1],
    [8, 8], [9, 8], [8, 9]
]

# 0 = No condition, 1 = Condition
y = [0, 0, 0, 1, 1, 1]

model = KNeighborsClassifier(n_neighbors=3)
model.fit(X, y)

patient = [[7, 8]]

prediction = model.predict(patient)

if prediction[0] == 1:
    print("Patient may have the condition")
else:
    print("Patient may not have the condition")

... Patient may have the condition
```

Question 25: Decision Tree for Iris Classification

```
[37] ✓ 13s X = iris.data
y = iris.target

model = DecisionTreeClassifier(random_state=1)
model.fit(X, y)

sepal_length = float(input("Sepal length: "))
sepal_width = float(input("Sepal width: "))
petal_length = float(input("Petal length: "))
petal_width = float(input("Petal width: "))

new_flower = [
    sepal_length, sepal_width,
    petal_length, petal_width
]

prediction = model.predict(new_flower)

print("Predicted species:", iris.target_names[prediction[0]])

... Sepal length: 5
Sepal width: 3
Petal length: 4
Petal width: 6
Predicted species: versicolor
```

Question 26: Linear Regression for Housing Price Prediction

```
import numpy as np
from sklearn.linear_model import LinearRegression

# Area in square feet
X = np.array([[1000], [1200], [1500], [1800], [2000]])
# House prices
y = np.array([200000, 240000, 300000, 360000, 400000])

model = LinearRegression()
model.fit(X, y)

area = float(input("Enter house area: "))

prediction = model.predict([[area]])

print("Predicted house price =", prediction[0])
```

... Enter house area: 100
Predicted house price = 19999.99999999995

Question 27: Logistic Regression for Customer Churn

```
[39] ✓ 13s
# Features: usage minutes, contract months
X = np.array([
    [100, 24], [120, 24], [80, 12],
    [300, 6], [280, 6], [350, 3]
])

# 0 = Not churned, 1 = Churned
y = np.array([0, 0, 0, 1, 1, 1])

model = LogisticRegression()
model.fit(X, y)

usage = float(input("Usage minutes: "))
contract = float(input("Contract duration: "))

prediction = model.predict([[usage, contract]])

if prediction[0] == 1:
    print("Customer may churn")
else:
    print("Customer may not churn")
```

... Usage minutes: 60
Contract duration: 5
Customer may not churn

Question 28: K-Means Customer Segmentation

```
[40]
✓ 0s
import numpy as np
from sklearn.cluster import KMeans

X = np.array([
    [1000, 2],
    [1200, 3],
    [1100, 2],
    [5000, 10],
    [5200, 11],
    [4800, 9]
])

model = KMeans(n_clusters=2, random_state=1, n_init=10)
model.fit(X)

new_customer = [[1500, 3]]
cluster = model.predict(new_customer)

print("Customer belongs to cluster:", cluster[0])

... Customer belongs to cluster: 0
```

Question 29: Accuracy, Precision, Recall and F1-Score

```
11]
✓ 0s
from sklearn.metrics import (
    accuracy_score, precision_score,
    recall_score, f1_score
)

# Actual and predicted values
y_test = [0, 1, 1, 0, 1, 0, 1, 1]
y_pred = [0, 1, 0, 0, 1, 0, 1, 0]

print("Accuracy =", accuracy_score(y_test, y_pred))
print("Precision =", precision_score(y_test, y_pred))
print("Recall =", recall_score(y_test, y_pred))
print("F1-score =", f1_score(y_test, y_pred))

... Accuracy = 0.75
Precision = 1.0
Recall = 0.6
F1-score = 0.75
```

Question 30: CART Regression for Used Car Price

```
import numpy as np
from sklearn.tree import DecisionTreeRegressor

# Features: mileage, age
X = np.array([
    [20000, 2], [30000, 3], [40000, 4],
    [60000, 5], [80000, 7], [100000, 9]
])

price = np.array([18000, 17000, 15000, 13000, 10000, 7000])

model = DecisionTreeRegressor(max_depth=3, random_state=1)
model.fit(X, price)

mileage = float(input("Mileage: "))
age = float(input("Age: "))

prediction = model.predict([[mileage, age]])

print("Predicted price =", prediction[0])
```

... Mileage: 10
Age: 40
Predicted price = 7000.0

Question 31: Customer Segmentation using K-Means



Question 32: Bivariate Analysis and Linear Regression for House Price



Question 33: Multiple Linear Regression for Car Price

```
data = pd.DataFrame({
    "Engine": [1.2, 1.5, 1.8, 2.0, 2.2, 2.5],
    "Horsepower": [80, 100, 120, 140, 160, 190],
    "Mileage": [20, 18, 16, 15, 13, 11],
    "Price": [8000, 10000, 13000, 16000, 19000, 23000]
})

X = data[["Engine", "Horsepower", "Mileage"]]
y = data["Price"]

model = LinearRegression()
model.fit(X, y)

predicted = model.predict(X)

print("R-squared =", r2_score(y, predicted))
print("Coefficients =", model.coef_)

new_car = [[1.8, 120, 16]]
print("Predicted price =", model.predict(new_car)[0])
```

... R-squared = 0.9983049479376866
Coefficients = [-2692.30769231 180.76923077 115.38461538]
Predicted price = 12961.538461538466
/usr/local/lib/python3.13/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but LinearRegression has feature names
warnings.warn(

Question 34: KNN Classification with Evaluation Metrics

```
[46]
✓ Os  ])
```

```
y = np.array([0, 0, 0, 0, 1, 1, 1, 1])

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=1
)

model = KNeighborsClassifier(n_neighbors=3)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Accuracy =", accuracy_score(y_test, y_pred))
print("Precision =", precision_score(y_test, y_pred))
print("Recall =", recall_score(y_test, y_pred))
print("F1-score =", f1_score(y_test, y_pred))

new_patient = [[30, 150]]
print("Prediction =", model.predict(new_patient))
```

```
... Accuracy = 1.0
Precision = 1.0
Recall = 1.0
F1-score = 1.0
Prediction = [0]
```

Question 35: K-Means Clustering for Retail Customers

```
[48]
✓ Os  )
```

```
"Spending": [500, 700, 600, 4000, 4500, 4200, 1500, 1700],
"Visits": [2, 3, 2, 10, 12, 11, 5, 6]
})

model = KMeans(n_clusters=3, random_state=1, n_init=10)


data["Cluster"] = model.fit_predict(
    data[["Spending", "Visits"]]
)

print(data)


print("Cluster centers:")
print(model.cluster_centers_)
```

```
...
   Spending  Visits  Cluster
0         500        2        1
1         700        3        1
2         600        2        1
3        4000       10        0
4        4500       12        0
5        4200       11        0
6        1500        5        2
7        1700        6        2
Cluster centers:
[[4.23333333e+03 1.10000000e+01]
 [6.00000000e+02 2.33333333e+00]
 [1.60000000e+03 5.50000000e+00]]
```

Question 36: Stock Price Variability from CSV



Question 37: Correlation between Study Time and Exam Score



Question 38: Temperature Analysis for Different Cities


```
[52] ✓ 0s
for i in range(3):
    print(cities[i])
    print("Mean =", mean[i])
    print("Standard deviation =", std[i])
    print("Range =", temperature_range[i])

highest_range = np.argmax(temperature_range)
most_consistent = np.argmin(std)

print("Highest temperature range:", cities[highest_range])
print("Most consistent city:", cities[most_consistent])


... Chennai
Mean = 27.2
Standard deviation = 1.7204650534085253
Range = 5
Bengaluru
Mean = 20.8
Standard deviation = 0.7483314773547882
Range = 2
Delhi
Mean = 32.2
Standard deviation = 1.7204650534085253
Range = 5
Highest temperature range: Chennai
Most consistent city: Bengaluru
```

Question 39: K-Means Clustering and Visualization



Question 40: Soccer Player Analysis using Pandas

...



Number of Players

Position

Position	Number of Players
Forward	3
Defender	2
Midfielder	1
Goalkeeper	1

```
players["Position"].value_counts().plot(kind="bar")
plt.title("Players by Position")
plt.xlabel("Position")
plt.ylabel("Number of Players")
plt.show()
```

...

Top 5 goal scorers:

Name	Goals
2 C	20
0 A	15
5 F	12
1 B	8
6 G	5

Top 5 salaries:

Name	Salary
2 C	7000
5 F	6000
0 A	5000
1 B	4500
4 E	4000

Average age = 25.571428571428573

Players above average age:

2 C
4 E
5 F

Name: Name, dtype: object